
GeoSpatialTools

Release 0.11.2

NOC Surface Processes

Mar 14, 2025

CONTENTS:

1	Introduction	1
1.1	GeoSpatialTools	1
1.2	Credits	1
1.2.1	Development Lead	1
1.2.2	Contributoring Developers	1
2	Installation	3
2.1	Via UV	3
2.1.1	Development mode	3
2.2	Via Pip	4
2.2.1	From Source	4
3	Bisection Search	5
3.1	Example	5
3.2	neighbours Module	6
4	Record Classes	7
4.1	Example	7
4.2	record Module	7
5	Shape Classes	9
5.1	Example	9
5.2	shape Module	10
6	Quadtree	13
6.1	Documentation	13
6.1.1	Inserting Records	13
6.1.2	Removing Records	13
6.1.3	Querying	13
6.2	Example	14
6.3	quadtree Module	14
7	OctTree	17
7.1	Documentation	17
7.1.1	Inserting Records	17
7.1.2	Removing Records	17
7.1.3	Querying	17
7.2	Example	18
7.3	octtree Module	19
8	K-D-Tree	21

8.1	Example	21
8.2	Documentation	22
8.2.1	Inserting Records	22
8.2.2	Removing Records	22
8.2.3	Querying	22
8.3	kdtree Module	22
9	Additional Modules	25
9.1	GreatCircle	25
9.2	Distance Metrics	27
9.3	Utils	28
	Python Module Index	29
	Index	31

INTRODUCTION

1.1 GeoSpatialTools

GeoSpatialTools is a [python3](#) library developed at [NOC](#) (National Oceanography Centre, Southampton, UK) for identifying neighbours in a geo-spatial context. This is designed to solve problems where one needs to identify data within a spatial range on the surface of the Earth. The library provides implementations of standard tools for neighbourhood searching, such as [k-d-tree](#) and [Quadtree](#) that have been adapted to account for spherical geometry, using a [haversine](#) distance metric.

The tool allows for spatial look-ups with $O(\log(n))$ complexity in time. Additionally, a simple 1-d nearest neighbours look-up is provided for sorted data using [bisection](#) search.

GeoSpatialTools also provides functionality for working with [great-circle](#) objects, for example intersecting great-circles.

1.2 Credits

1.2.1 Development Lead

- Joseph T. Siddons <josidd@noc.ac.uk> @josidd

1.2.2 Contributing Developers

- Richard C. Cornes <rcornes@noc.ac.uk> @ricorne

INSTALLATION

GeoSpatialTools is not currently available on PyPI so must be installed either from source or directly from the GitLab repository. Versions of python between 3.9 and 3.13, inclusive, are supported, however the recommended version of python is 3.12.

We recommend the installation of GeoSpatialTools using the `uv` package manager, however it can be installed using `pip`.

The only required dependency of the project is `NumPy`. Additional dependency `polars` is required to run the `Jupyter` notebooks.

2.1 Via UV

You can install the library directly from the GitLab repository, adding the library to your current `uv` virtual environment. This will add the library as a dependency in your current project.

```
uv add git+ssh://git@git.noc.ac.uk/nocsurfaceprocesses/geospatialtools.git
```

2.1.1 Development mode

If you wish to contribute to GeoSpatialTools you can install the library in development mode. This will require cloning the repository and creating a new `uv` environment.

```
# Get the code
git clone git@git.noc.ac.uk/nocsurfaceprocesses/geospatialtools.git
cd geospatialtools

# Install with all dependencies and create an environment with python 3.12
uv sync --all-extras --dev --python 3.12

# Load the environment
source .venv/bin/activate

# Run the unit tests
uv run pytest test
```

Note

The recommended python version is python 3.12. By default, `uv` creates a virtual environment in `.venv`.

2.2 Via Pip

The library can be installed via pip with the following command:

```
pip install git+ssh://git@git.noc.ac.uk/nocsurfaceprocesses/geospatialtools.git
```

2.2.1 From Source

Alternatively, you can clone the repository and install using pip (or conda if preferred). This installs in editable mode.

```
git clone git@git.noc.ac.uk/nocsurfaceprocesses/geospatialtools.git
cd geospatialtools
python -m venv venv
source venv/bin/activate
pip install -e .
```


BISECTION SEARCH

Bisection can be used to find the nearest neighbour in a sorted one-dimensional list of search values in $O(\log(n))$ time complexity.

The implementation in *GeoSpatialTools* makes use of the *bisect* library, which is part of the Python standard library. The input types are numeric types, which can include `int`, `float`, or `datetime.datetime` values.

The bisection approach repeatedly splits the list of search values in two at the mid-index. The query value is compared to the search value at the mid-index. If the query value is larger than the search value at the mid-index, then the search values after the mid-index become the new search values. If the query value is smaller than the search value at the mid-index then the search values before the mid-index become the new search values. This bisecting is repeated (successively halving the number of search values) until one values remain. The nearest neighbour is either the value at the remaining index, or the value at the index one above.

Note

The above assumes that the list of search values is sorted in increasing order. The opposite applies if the list is sorted in reverse.

Warning

The input values must be sorted

3.1 Example

```
from GeoSpatialTools import find_nearest
import numpy as np

search_values: list[float] = list(np.random.randn(50))
search_values.sort()

query_value: float = 0.45
neighbour_index: int = find_nearest(
    vals=search_values,
    test=query_value,
)
neighbour_value: float = search_values[neighbour_index]
```

3.2 neighbours Module

Functions for finding nearest neighbours using bisection. Nearest neighbours can be found with $O(\log(n))$ time-complexity.

Data for these functions must be sorted, otherwise incorrect values may be returned.

exception `GeoSpatialTools.neighbours.SortedError`

Error class for Sortedness

exception `GeoSpatialTools.neighbours.SortedWarning`

Warning class for Sortedness

`GeoSpatialTools.neighbours.find_nearest(vals, test, check_sorted=True)`

Find the nearest value in a list of values for each test value.

Uses bisection for speediness!

Returns a list containing the index of the nearest neighbour in vals for each value in test. Or the index of the nearest neighbour if test is a single value.

Parameters

- **vals** (*list[Numeric]*) – List of values - this is the pool of values for which we are looking for a nearest match. This list **MUST** be sorted. Sortedness is not checked, nor is the list sorted.
- **test** (*Numeric | list[Numeric]*) – Query value(s)
- **check_sorted** (*bool*) – Optionally check that the input vals is sorted. Raises an error if set to True (default), displays a warning if set to False.

Returns

Index, or list of indices, of nearest value, or values.

Return type

`int | list[int]`

RECORD CLASSES

Record classes in `GeoSpatialTools` form the back-bone of the data structures within the library. They represent a consistent input data-type across all classes in the library.

There are two classes of Record:

- `GeoSpatialTools.record.Record` for two-dimensional data structures defined by `lon` (longitude) and `lat` (latitude). Optionally, one can pass `datetime` and `uid`, as well as additional data attributes with keyword arguments.
- `GeoSpatialTools.record.SpaceTimeRecord` for three-dimensional data structures defined by `lon` (longitude), `lat` (latitude), and `datetime`. Optionally, one can pass `uid`, as well as additional data attributes with keyword arguments.

Only the positional, `datetime`, and `uid` attributes are used for equality tests. `Record` objects are used for `QuadTree` and `KDTree` objects, whereas `SpaceTimeRecord` objects must be used for `OctTree`.

Note

`Record` and `SpaceTimeRecord` are exposed at the `GeoSpatialTools` level.

4.1 Example

```
from GeoSpatialTools import Record

record: Record = Record(lon=-151.2, lat=42.7, uid="foo")
dist: float = record.distance(Record(-71.1, -23.2, uid="bar"))
```

4.2 record Module

Record objects used for containing data passed to `QuadTree`, `OctTree` and `KDTree` classes. Require positions defined by “lon” and “lat”, `SpaceTimeRecord` objects also require “datetime”. Optional fields are “uid”, other data can be passed as keyword arguments. Only positional, temporal, and uid values are used for equality checks.

Distances between records is calculated using Haversine distance.

Classes prefixed by “SpaceTime” include a temporal dimension and should be used with `OctTree` classes.

class `GeoSpatialTools.record.Record(lon, lat, datetime=None, uid=None, fix_lon=True, **data)`

Record class

This is a simple instance of an record, it requires position data. It can optionally include `datetime`, a `UID`, and extra data passed as keyword arguments.

Equality is checked only on the required fields + UID if it is specified.

Parameters

- **lon** (*float*) – Horizontal coordinate
- **lat** (*float*) – Vertical coordinate
- **datetime** (*datetime* / *None*) – Datetime of the record
- **uid** (*str* / *None*) – Unique Identifier
- **fix_lon** (*bool*) – Force longitude to -180, 180
- ****data** – Additional data passed to the Record for use by other functions or classes.

distance(*other*)

Compute the Haversine distance to another Record

Return type

float

class GeoSpatialTools.record.SpaceTimeRecord(*lon, lat, datetime, uid=None, fix_lon=True, **data*)

ICOADS Record class.

This is a simple instance of an ICOARDS record, it requires position and temporal data. It can optionally include a UID and extra data.

The temporal component was designed to use *datetime* values, however all methods will work with numeric datetime information - for example a pentad, timestamp, julian day, etc. Note that any uses within an OctTree and SpaceTimeRectangle must also have timedelta values replaced with numeric ranges in this case.

Equality is checked only on the required fields + UID if it is specified.

Parameters

- **lon** (*float*) – Horizontal coordinate (longitude).
- **lat** (*float*) – Vertical coordinate (latitude).
- **datetime** (*datetime.datetime*) – Datetime of the record. Can also be a numeric value such as pentad. Comparisons between Records with datetime and Records with numeric datetime will fail.
- **uid** (*str* / *None*) – Unique Identifier.
- **fix_lon** (*bool*) – Force longitude to -180, 180
- ****data** – Additional data passed to the SpaceTimeRecord for use by other functions or classes.

distance(*other*)

Compute the Haversine distance to another SpaceTimeRecord. Only computes spatial distance.

Return type

float

SHAPE CLASSES

The `GeoSpatialTools.shape` module defines various classes that can be used to define the boundary for `QuadTree` and `OctTree` classes, or query regions for the same.

`Rectangle` and `SpaceTimeRectangle` classes are used to define the boundaries for `QuadTree` and `OctTree` classes respectively. They are defined by the bounding box in space (and time for a `SpaceTimeRectangle`).

`Ellipse` and `SpaceTimeEllipse` classes are defined by `lon` and `lat` indicating the centre of the ellipse, `a` and `b` indicating the length of the semi-major and semi-minor axes respectively, and `theta` indicating the angle of the ellipse. `SpaceTimeEllipse` classes also require `start` and `end` datetime values. The `SpaceTimeEllipse` is an elliptical cylinder where the height is represented by the time dimension.

`Rectangle` and `Ellipse` classes can be used to define a query shape for a `QuadTree`, using `QuadTree.query` and `QuadTree.query_ellipse` respectively.

`SpaceTimeRectangle` and `SpaceTimeEllipse` classes can be used to define a query shape for a `OctTree`, using `OctTree.query` and `OctTree.query_ellipse` respectively.

5.1 Example

```
from GeoSpatialTools.shape import Rectangle, SpaceTimeEllipse
from datetime import datetime
from math import pi

rectangle: Rectangle = Rectangle(
    west=-180,
    east=180,
    south=-90,
    north=90,
)

ellipse: SpaceTimeEllipse = SpaceTimeEllipse(
    lon=23.4,
    lat=-17.9,
    a=103,
    b=71,
    theta=pi/3,
    start=datetime(2009, 2, 13, 19, 30),
    end=datetime(2010, 7, 2, 3, 45),
)
```

5.2 shape Module

Shape objects for defining QuadTree or OctTree classes, or for defining a query region for QuadTree and OctTree classes.

Distances between shapes, or between shapes and Records uses the Haversine distance.

All shape objects account for spherical geometry and the wrapping of longitude at -180, 180 degrees.

Classes prefixed by “SpaceTime” include a temporal dimension and should be used with OctTree classes.

class GeoSpatialTools.shape.**Ellipse**(*lon, lat, a, b, theta*)

A simple Ellipse Class for an ellipse on the surface of a sphere.

Parameters

- **lon** (*float*) – Horizontal centre of the Ellipse
- **lat** (*float*) – Vertical centre of the Ellipse
- **a** (*float*) – Length of the semi-major axis
- **b** (*float*) – Length of the semi-minor axis
- **theta** (*float*) – Angle of the semi-major axis from horizontal anti-clockwise in radians

contains(*point*)

Test if a Record is contained within the Ellipse

Return type

bool

nearby_rect(*rect*)

Test if a Rectangle is near to the Ellipse

Return type

bool

class GeoSpatialTools.shape.**Rectangle**(*west, east, south, north*)

A simple Rectangle class for GeoSpatial analysis. Defined by a bounding box.

Parameters

- **west** (*float*) – Western boundary of the Rectangle
- **east** (*float*) – Eastern boundary of the Rectangle
- **south** (*float*) – Southern boundary of the Rectangle
- **north** (*float*) – Northern boundary of the Rectangle

contains(*point*)

Test if a Record is contained within the Rectangle

Return type

bool

property edge_dist: float

Approximate maximum distance from the centre to an edge

intersects(*other*)

Test if another Rectangle object intersects this Rectangle

Return type

bool

property lat: float

Centre latitude of the Rectangle

property lat_range: float

Latitude range of the Rectangle

property lon: float

Centre longitude of the Rectangle

property lon_range: float

Longitude range of the Rectangle

nearby(*point, dist*)

Check if Record is nearby the Rectangle

Return type

bool

class GeoSpatialTools.shape.SpaceTimeEllipse(*lon, lat, a, b, theta, start, end*)

A simple SpaceTimeEllipse Class for an ellipse on the surface of a sphere with an additional time dimension.

The representation of the shape is an elliptical cylinder, with the time dimension representing the height of the cylinder.

Parameters

- **lon** (*float*) – Horizontal centre of the SpaceTimeEllipse
- **lat** (*float*) – Vertical centre of the SpaceTimeEllipse
- **a** (*float*) – Length of the semi-major axis
- **b** (*float*) – Length of the semi-minor axis
- **theta** (*float*) – Angle of the semi-major axis from horizontal anti-clockwise in radians
- **start** (*datetime.datetime*) – Start date of the SpaceTimeEllipse
- **end** (*datetime.datetime*) – Send date of the SpaceTimeEllipse

contains(*point*)

Test if a SpaceTimeRecord is contained within the SpaceTimeEllipse

Return type

bool

nearby_rect(*rect*)

Test if a SpaceTimeRectangle is near to the SpaceTimeEllipse

Return type

bool

class GeoSpatialTools.shape.SpaceTimeRectangle(*west, east, south, north, start, end*)

A simple SpaceTimeRectangle class for GeoSpatial analysis. Defined by a bounding box in space and time.

Parameters

- **west** (*float*) – Western boundary of the SpaceTimeRectangle
- **east** (*float*) – Eastern boundary of the SpaceTimeRectangle
- **south** (*float*) – Southern boundary of the SpaceTimeRectangle
- **north** (*float*) – Northern boundary of the SpaceTimeRectangle

- **start** (*datetime.datetime*) – Start datetime of the SpaceTimeRectangle
- **end** (*datetime.datetime*) – End datetime of the SpaceTimeRectangle

property centre_datetime: datetime

The midpoint time of the SpaceTimeRectangle

contains(*point*)

Test if a SpaceTimeRecord is contained within the SpaceTimeRectangle

Return type

bool

property edge_dist: float

Approximate maximum distance from the centre to an edge

intersects(*other*)

Test if another SpaceTimeRectangle object intersects this SpaceTimeRectangle.

Return type

bool

property lat: float

Centre latitude of the SpaceTimeRectangle

property lat_range: float

Latitude range of the SpaceTimeRectangle

property lon: float

Centre longitude of the SpaceTimeRectangle

property lon_range: float

Longitude range of the SpaceTimeRectangle

nearby(*point, dist, t_dist*)

Check if SpaceTimeRecord is nearby the SpaceTimeRectangle

Determines if a SpaceTimeRecord that falls on the surface of Earth is nearby to the rectangle in space and time. This calculation uses the Haversine distance metric.

Distance from rectangle to point is challenging on the surface of a sphere, this calculation will return false positives as a check based on the distance from the centre of the rectangle to the corners, or to its Eastern edge (if the rectangle crosses the equator) is used in combination with the input distance.

The primary use-case of this method is for querying an OctTree for nearby SpaceTimeRecords.

Parameters

- **point** (*SpaceTimeRecord*)
- **dist** (*float*,)
- **t_dist** (*datetime.timedelta*)

Returns

bool

Return type

True if the point is $\leq \text{dist} + \max(\text{dist}(\text{centre}, \text{corners}))$

property time_range: timedelta

The time extent of the Rectangle

QUADTREE

A Quadtree is a data-structure where each internal node has exactly four children, and are used to recursively partition a two-dimensional spatial domain. Each child node is itself a Quadtree, whose spatial domain represents one of the quadrants (north-west, north-east, south-west, south-east) of its parent's domain. The partitioning of data in this way is dependent on the spatial density of data inserted into the Quadtree. The Quadtree is typically initialised with a capacity value, once the capacity is reached (by inserting data points), the Quadtree divides and subsequent data points are added to the appropriate child-node.

Quadtree structures allow for fast identification of data within some query region. The structure of the tree ensures that only nodes whose domain boundary intersects (or contains or is contained by) the query region are evaluated. The time-complexity of these query operations is $O(\log(n))$, the space-complexity of a Quadtree is $O(n)$.

Typically, it is assumed that the data uses a cartesian coordinate system, so comparisons between boundaries and query shapes utilise cartesian geometry and euclidean distances. The implementation of Quadtree within this library, the `QuadTree` class, utilises the Haversine distance as a metric for identifying records within the queried region. This allows the Quadtree to account for the spherical geometry of the Earth. Boundary checks with query regions also account for the wrapping of longitude at -180, 180 degrees.

The `QuadTree` object is defined by a bounding box, i.e. boundaries at the western, eastern, southern, and northern edges of the data that will be inserted into the `QuadTree`. Additionally, a capacity and maximum depth can be provided. If the capacity is exceeded whilst inserting records the `QuadTree` will divide and new records will be inserted into the appropriate child `QuadTree`. The maximum depth is the maximum height of the `QuadTree`, if capacity is also specified then this will be overridden if the `QuadTree` is at this depth, and the `QuadTree` will not divide.

6.1 Documentation

6.1.1 Inserting Records

A `Record` can be added to an `QuadTree` with `QuadTree.insert` which will return `True` if the operation was successful, `False` otherwise. The `QuadTree` is modified in place.

6.1.2 Removing Records

A `Record` can be removed from an `QuadTree` with `QuadTree.remove` which will return `True` if the operation was successful, `False` otherwise. The `QuadTree` is modified in place.

6.1.3 Querying

The `QuadTree` class defined in `GeoSpatialTools.quadtree` can be queried in the following ways:

- with a `Record`, a spatial range with `QuadTree.nearby_points`. All points within the spatial range of the `Record` will be returned in a list.

- with a Rectangle using `QuadTree.query`. All points within the specified Rectangle will be returned in a list.
- with a Ellipse using `QuadTree.query_ellipse`. All points within the specified Ellipse will be returned in a list.

6.2 Example

```
from GeoSpatialTools import QuadTree, Record, Rectangle
from random import choice

lon_range = list(range(-180, 180))
lat_range = list(range(-90, 90))

N_samples = 1000

# Construct Tree
boundary = Rectangle(
    west=-180,
    east=180,
    south=-90,
    north=90,
) # Full domain
quadtree = QuadTree(boundary)

# Populate the tree
records: list[Record] = [
    Record(
        choice(lon_range),
        choice(lat_range),
    ) for _ in range(N_samples)
]
for record in records:
    quadtree.insert(record)

dist: float = 340 # km

# Find all Records that are 340km away from test_value
neighbours: list[Record] = quadtree.nearby_points(test_value, dist)
```

6.3 quadtree Module

Constructors for QuadTree classes that can decrease the number of comparisons for detecting nearby records for example. This is an implementation that uses Haversine distances for comparisons between records for identification of neighbours.

class `GeoSpatialTools.quadtree.QuadTree`(*boundary, capacity=5, depth=0, max_depth=None*)

A Simple QuadTree class for PyCOADS

Parameters

- **boundary** (`Rectangle`) – The bounding Rectangle of the QuadTree

- **capacity** (*int*) – The capacity of each cell, if `max_depth` is set then a cell at the maximum depth may contain more points than the capacity.
- **depth** (*int*) – The current depth of the cell. Initialises to zero if unset.
- **max_depth** (*int* / *None*) – The maximum depth of the QuadTree. If set, this can override the capacity for cells at the maximum depth.

divide()

Divide the QuadTree

insert(point)

Insert a point into the QuadTree

Return type

`bool`

len(_current_len=0)

Get the number of points in the QuadTree

Return type

`int`

nearby_points(point, dist, points=None)

Get all Records contained in the QuadTree that are nearby another query Record.

Query the QuadTree to find all Records within the QuadTree that are nearby to the query Record. This search should be faster than searching through all records, since only QuadTree children whose boundaries are close to the query Record are evaluated.

Parameters

- **point** (*Record*) – The query point.
- **dist** (*float*) – The distance for comparison. Note that Haversine distance is used as the distance metric as the query Record and QuadTree are assumed to lie on the surface of Earth.
- **points** (*Records* / *None*) – List of Records already found. Most use cases will be to not set this value, since it's main use is for passing onto the children QuadTrees.

Returns

A list of Records whose distance to the query Record is \leq dist, and the datetimes of the Records fall within the datetime range of the query Record.

Return type

`list[Record]`

query(rect, points=None)

Get Records contained within the QuadTree that fall in a Rectangle

Parameters

rect (*Rectangle*)

Returns

The Record values contained within the QuadTree that fall within the bounds of rect.

Return type

`list[Record]`

query_ellipse(ellipse, points=None)

Get Records contained within the QuadTree that fall in a Ellipse

Parameters

ellipse ([Ellipse](#))

Returns

The Record values contained within the QuadTree that fall within the bounds of ellipse.

Return type

list[[Record](#)]

remove(*point*)

Remove a Record from the QuadTree if it is in the QuadTree.

Returns True if the Record is removed.

Return type

bool

OCTTREE

An Octtree is an extension of the Quadtree into a third dimension. In standard Octtree implementations the third dimension is treated as another spatial dimension, in that distance checks are performed using Euclidean distances. Here, the third dimension is considered to be a time dimension. Any look-ups using the Octtree require a `timedelta` to be provided, so that any records falling within the spatial range are returned only if they also fall within the time range defined by the `timedelta`.

Whilst the Quadtree divides into 4 children after the capacity is reached, the Octtree divides into 8 children. The divisions are at the longitude midpoint, the latitude midpoint, and the datetime midpoint of the boundary.

The implementation of Octtree within this library, the `OctTree` class, utilises the Haversine distance as a metric for identifying records within the queried region. This allows the Octtree to account for the spherical geometry of the Earth. Boundary checks with query regions also account for the wrapping of longitude at -180, 180 degrees.

The `OctTree` object is defined by a bounding box in space and time, i.e. boundaries at the western, eastern, southern, and northern edges as well as the start and end datetimes of the data that will be inserted into the `OctTree`. Additionally, a capacity and maximum depth can be provided. If the capacity is exceeded whilst inserting records the `OctTree` will divide and new records will be inserted into the appropriate child `OctTree`. The maximum depth is the maximum height of the `OctTree`, if capacity is also specified then this will be overridden if the `OctTree` is at this depth, and the `OctTree` will not divide.

7.1 Documentation

7.1.1 Inserting Records

A `SpaceTimeRecord` can be added to an `OctTree` with `OctTree.insert` which will return `True` if the operation was successful, `False` otherwise. The `OctTree` is modified in place. Records that fall outside the bounds of the `OctTree` will not be inserted as the boundary is fixed.

7.1.2 Removing Records

A `SpaceTimeRecord` can be removed from an `OctTree` with `OctTree.remove` which will return `True` if the operation was successful, `False` otherwise. The `OctTree` is modified in place.

7.1.3 Querying

The `OctTree` class defined in `GeoSpatialTools.octtree` can be queried in the following ways:

- with a `SpaceTimeRecord`, a spatial range, and a time range (specified by a `datetime.timedelta`) with `OctTree.nearby_points`. All points within the spatial range and time range of the `SpaceTimeRecord` will be returned in a list.
- with a `SpaceTimeRectangle` using `OctTree.query`. All points within the specified `SpaceTimeRectangle` will be returned in a list.

- with a SpaceTimeEllipse using OctTree.query_ellipse. All points within the specified SpaceTimeEllipse will be returned in a list.

7.2 Example

```

from GeoSpatialTools import OctTree, SpaceTimeRecord, SpaceTimeRectangle
from datetime import datetime, timedelta
from random import choice
from polars import datetime_range

lon_range = list(range(-180, 180))
lat_range = list(range(-90, 90))

dates = datetime_range(
    start=datetime(2009, 1, 1, 0, 0),
    end=datetime(2009, 2, 1, 0, 0),
    interval=timedelta(hours=1),
    closed="left",
    eager=True,
)
N_samples = 1000

# Construct Tree
boundary = SpaceTimeRectangle(
    west=-180,
    east=180,
    south=-90,
    north=90,
    start=datetime(2009, 1, 1, 0),
    end=datetime(2009, 1, 2, 23),
) # Full domain
octtree = OctTree(boundary)

# Populate the tree
records: list[SpaceTimeRecord] = [
    SpaceTimeRecord(
        choice(lon_range),
        choice(lat_range),
        choice(dates)
    ) for _ in range(N_samples)
]
for record in records:
    octtree.insert(record)

test_value: SpaceTimeRecord = SpaceTimeRecord(
    lon=47.6, lat=-31.1, datetime=datetime(2009, 1, 23, 17, 41)
)
dist: float = 340 # km
t_dist = timedelta(hours=4)

# Find all Records that are 340km away from test_value, and within 4 hours
# of test_value

```

(continues on next page)

(continued from previous page)

```

neighbours: list[SpaceTimeRecord] = octtree.nearby_points(
    test_value, dist, t_dist
)

```

7.3 octtree Module

Constructors for OctTree classes that can decrease the number of comparisons for detecting nearby records for example. This is an implementation that uses Haversine distances for comparisons between records for identification of neighbours.

class GeoSpatialTools.octtree.**OctTree**(*boundary, capacity=5, depth=0, max_depth=None*)

A Simple OctTree class for PyCOADS.

Acts as a space-time OctTree on the surface of Earth, allowing for querying nearby points faster than searching a full DataFrame. As SpaceTimeRecords are added to the OctTree, the OctTree divides into 8 children as the capacity is reached. Additional SpaceTimeRecords are then added to the children where they fall within the child OctTree's boundary.

SpaceTimeRecords already part of the OctTree before divided are not distributed to the children OctTrees.

Whilst the OctTree has a temporal component, and was designed to utilise datetime / timedelta objects, numeric values and ranges can be used. This usage must be consistent for the boundary and all SpaceTimeRecords that are part of the OctTree. This allows for usage of pentad, timestamp, Julian day, etc. as datetime values.

Parameters

- **boundary** (*SpaceTimeRectangle*) – The bounding SpaceTimeRectangle of the QuadTree
- **capacity** (*int*) – The capacity of each cell, if *max_depth* is set then a cell at the maximum depth may contain more points than the capacity.
- **depth** (*int*) – The current depth of the cell. Initialises to zero if unset.
- **max_depth** (*int | None*) – The maximum depth of the QuadTree. If set, this can override the capacity for cells at the maximum depth.

divide()

Divide the QuadTree

insert(*point*)

Insert a SpaceTimeRecord into the QuadTree.

Note that the SpaceTimeRecord can have numeric datetime values if that is consistent with the OctTree.

Return type

bool

len(*_current_len=0*)

Get the number of points in the OctTree

Return type

int

nearby_points(*point, dist, t_dist, points=None*)

Get all SpaceTimeRecords contained in the OctTree that are nearby another query SpaceTimeRecord.

Query the OctTree to find all SpaceTimeRecords within the OctTree that are nearby to the query SpaceTimeRecord. This search should be faster than searching through all records, since only OctTree children whose boundaries are close to the query SpaceTimeRecord are evaluated.

Parameters

- **point** (*SpaceTimeRecord*) – The query point.
- **dist** (*float*) – The distance for comparison. Note that Haversine distance is used as the distance metric as the query *SpaceTimeRecord* and *OctTree* are assumed to lie on the surface of Earth.
- **t_dist** (*datetime.timedelta*) – Max time gap between *SpaceTimeRecords* within the *OctTree* and the query *SpaceTimeRecord*. Can be numeric if the *OctTree* boundaries, *SpaceTimeRecords*, and query *SpaceTimeRecord* have numeric datetime values and ranges.
- **points** (*List[SpaceTimeRecord] / None*) – List of *SpaceTimeRecords* already found. Most use cases will be to not set this value, since it's main use is for passing onto the children *OctTrees*.

Returns

A list of *SpaceTimeRecords* whose distance to the query *SpaceTimeRecord* is \leq dist, and the datetimes of the *SpaceTimeRecords* fall within the datetime range of the query *SpaceTimeRecord*.

Return type

list[SpaceTimeRecord]

query(*rect*, *points=None*)

Get *SpaceTimeRecords* contained within the *OctTree* that fall in a *SpaceTimeRectangle*

Parameters

rect (*SpaceTimeRectangle*)

Returns

The *SpaceTimeRecord* values contained within the *OctTree* that fall within the bounds of *rect*.

Return type

List[SpaceTimeRecord]

query_ellipse(*ellipse*, *points=None*)

Get *SpaceTimeRecords* contained within the *OctTree* that fall in a *SpaceTimeEllipse*

Parameters

ellipse (*SpaceTimeEllipse*)

Returns

The *SpaceTimeRecord* values contained within the *OctTree* that fall within the bounds of *ellipse*.

Return type

List[SpaceTimeRecord]

remove(*point*)

Remove a *SpaceTimeRecord* from the *OctTree* if it is in the *OctTree*.

Returns True if the *SpaceTimeRecord* is removed.

Return type

bool

K-D-TREE

A K-D-Tree is a data structure that operates in a similar way to bisection or a binary tree, and can be used to find the nearest neighbour. For k -dimensional data (i.e. the data has k features), a binary tree is constructed by bisecting the data along each of the k dimensions in sequence. The first layer bisects the data along the first dimension, the second layer bisects each of the previous bisection results along the 2nd dimension (the data is now partitioned into 4), and so on. The pattern repeats after the k -th layer, until a single point of data remains in each leaf node. A K-D-Tree that bisects data and results in each leaf node containing a single value is called referred to as a balanced K-T-Tree.

To find the data point that is closest to a point in the tree, one descends the tree comparing the query point to the partition value in each dimension. The final leaf node should be the closest point, however there may be a point closer if the query point is close to a previous partition value, so some back tracking is performed to either confirm, or update, the closest point.

A K-D-Tree can typically find the nearest neighbour in $O(\log(n))$ time complexity, and the data structure has $O(n)$ space-complexity.

Most implementations of K-D-Tree assume that the coordinates use a cartesian geometry and therefore use a simple Euclidean distance to identify the nearest neighbour. The implementation in `GeoSpatialTools.kdtree` assumes a spherical geometry on the surface of the Earth and uses the Haversine distance to identify neighbours. The implementation has been designed to account for longitude wrapping at -180, 180 degrees. The `GeoSpatialTools.kdtree.KDTree` class is a 2-D-Tree, the dimensions are longitude and latitude. The object is initialised with data in the form of a list of `GeoSpatialTools.quadtrees.Record` objects. A maximum depth value (`max_depth`) can be provided, if this is set then the partitioning will stop after `max_depth` partitionings, the leaf nodes may contain more than one `Record`.

8.1 Example

```
from GeoSpatialTools import KDTree, Record
from random import choice

lon_range = list(range(-180, 180))
lat_range = list(range(-90, 90))
N_samples = 1000

records: list[Record] = [Record(choice(lon_range), choice(lat_range)) for _ in range(N_
↪samples)]
# Construct Tree
kdtree = KDTree(records)

test_value: Record = Record(lon=47.6, lat=-31.1)
neighbours, dist = kdtree.query(test_value)
```

8.2 Documentation

Note

Insertion and deletion operations may cause the KDTree to become un-balanced.

8.2.1 Inserting Records

A Record can be inserted in to a KDTree with the `KDTree.insert` method. The method will return `True` if the Record was inserted into the KDTree, `False` otherwise. A Record will not be added if it is already contained within the KDTree, to add the Record anyway use the `KDTree._insert` method.

8.2.2 Removing Records

A Record can be removed from a KDTree with the `KDTree.delete` method. The method will return `True` if the Record was successfully removed, `False` otherwise (for example if the Record is not contained within the KDTree).

8.2.3 Querying

The nearest neighbour Record contained within a KDTree to a query Record can be found with the `KDTree.query` method. This will return a tuple containing the list of Record objects from the KDTree with minimum distance to the query Record, and the minimum distance.

8.3 kdtree Module

An implementation of KDTree using Haversine Distance for GeoSpatial analysis. Useful tool for quickly searching for nearest neighbours. The implementation is a K=2 or 2DTree as only 2 dimensions (longitude and latitude) are used.

Haversine distances are used for comparisons, so that the spherical geometry of the earth is accounted for.

class `GeoSpatialTools.kdtree.KDTree(points, depth=0, max_depth=20)`

A Haverine distance implementation of a balanced KDTree.

This implementation is a `_balanced_` KDTree, each leaf node should have the same number of points (or differ by 1 depending on the number of points the KDTree is initialised with).

The KDTree partitions in each of the lon and lat dimensions alternatively in sequence by splitting at the median of the dimension of the points assigned to the branch.

Parameters

- **points** (*list[Record]*) – A list of `GeoSpatialTools.Record` instances.
- **depth** (*int*) – The current depth of the KDTree, you should set this to 0, it is used internally.
- **max_depth** (*int*) – The maximum depth of the KDTree. The leaf nodes will have depth no larger than this value. Leaf nodes will not be created if there is only 1 point in the branch.

delete(point)

Delete a Record from the KDTree. May unbalance the KDTree

Return type

`bool`

insert(point)

Insert a Record into the KDTree. May unbalance the KDTree.

The point will not be inserted if it is already in the KDTree.

Return type
bool

query(*point*)

Find the nearest Record within the KDTree to a query Record

Return type
Tuple[List[Record], float]

ADDITIONAL MODULES

9.1 GreatCircle

Constructors and methods for interacting with GreatCircle objects, including comparisons between GreatCircle objects.

class `GeoSpatialTools.great_circle.GreatCircle(lon0, lat0, lon1, lat1, R=6371)`

A GreatCircle object for a pair of positions.

Construct a great circle path between a pair of positions.

<https://www.boeing-727.com/Data/fly%20odds/distance.html>

Parameters

- **lon0** (*float*) – Longitude of start position.
- **lat0** (*float*) – Latitude of start position.
- **lon1** (*float*) – Longitude of end position.
- **lat1** (*float*) – Latitude of end position.
- **R** (*float*) – Radius of the sphere. Default is Earth radius in km (6371.0).

dist_from_point(*lon, lat*)

Compute distance from the GreatCircle to a point on the sphere.

Parameters

- **lon** (*float*) – Longitude of the position to test.
- **lat** (*float*) – Latitude of the position to test.

Returns

Minimum distance between point and the GreatCircle arc.

Return type

float

intersection(*other, epsilon=0.01*)

Determine intersection position with another GreatCircle.

Determine the location at which the GreatCircle intersects another GreatCircle arc. (To within some epsilon threshold).

Returns *None* if there is no solution - either because there is no intersection point, or the planes generated from the arc and centre of the sphere are identical.

Parameters

- **other** (*GreatCircle*) – Intersecting GreatCircle object
- **epsilon** (*float*) – Threshold for intersection

Returns

Position of intersection

Return type

(float, float) | None

intersection_angle(*other, epsilon=0.01*)

Get angle of intersection with another GreatCircle.

Get the angle of intersection with another GreatCircle arc. Returns None if there is no intersection.

The intersection angle is computed using the normals of the planes formed by the two intersecting great circle objects.

Parameters

- **other** (*GreatCircle*) – Intersecting GreatCircle object
- **epsilon** (*float*) – Threshold for intersection

Returns

Intersection angle in degrees

Return type

float | None

GeoSpatialTools.great_circle.cartesian_to_lonlat(*x, y, z, to_radians=False*)

Get lon, and lat from cartesian coordinates.

Parameters

- **x** (*float*) – x coordinate
- **y** (*float*) – y coordinate
- **z** (*float*) – z coordinate
- **to_radians** (*bool*) – Return angles in radians. Otherwise return values in degrees.

Return type

Tuple[float, float]

Returns

- (*float, float*)
- *lon, lat*

GeoSpatialTools.great_circle.polar_to_cartesian(*lon, lat, R=6371, to_radians=True, normalised=True*)

Convert from polars coordinates to cartesian.

Get cartesian coordinates from spherical polar coordinates. Default behaviour assumes lon and lat, so converts to radians. Set *to_radians=False* if the coordinates are already in radians.

Parameters

- **lon** (*float*) – Longitude.
- **lat** (*float*) – Latitude.
- **R** (*float*) – Radius of sphere.

- **to_radians** (*bool*) – Convert lon and lat to radians.
- **normalised** (*bool*) – Return normalised vector (ignore R value).

Returns

x, y, z cartesian coordinates.

Return type

(float, float, float)

9.2 Distance Metrics

Functions for computing navigational information. Can be used to add navigational information to DataFrames.

`GeoSpatialTools.distance_metrics.bearing(lon0, lat0, lon1, lat1)`

Compute the bearing of a track from (lon0, lat0) to (lon1, lat1).

Duplicated from geo-py

Parameters

- **lon0** (*float*,) – Longitude of start point
- **lat0** (*float*,) – Latitude of start point
- **lon1** (*float*,) – Longitude of target point
- **lat1** (*float*,) – Latitude of target point

Returns

bearing – The bearing from point (lon0, lat0) to point (lon1, lat1) in degrees.

Return type

float

`GeoSpatialTools.distance_metrics.destination(lon, lat, bearing, distance)`

Compute destination of a great circle path.

Compute the destination of a track started from 'lon', 'lat', with 'bearing'. Distance is in units of km.

Duplicated from geo-py

Parameters

- **lon** (*float*) – Longitude of initial position
- **lat** (*float*) – Latitude of initial position
- **bearing** (*float*) – Direction of track
- **distance** (*float*) – Distance to travel

Returns

destination – Longitude and Latitude of final position

Return type

tuple[float, float]

`GeoSpatialTools.distance_metrics.gcd_slc(lon0, lat0, lon1, lat1)`

Compute great circle distance on earth surface between two locations.

Parameters

- **lon0** (*float*) – Longitude of position 0
- **lat0** (*float*) – Latitude of position 0

- **lon1** (*float*) – Longitude of position 1
- **lat1** (*float*) – Latitude of position 1

Returns

dist – Great circle distance between position 0 and position 1.

Return type

float

`GeoSpatialTools.distance_metrics.haversine(lon0, lat0, lon1, lat1)`

Compute Haversine distance between two points.

Parameters

- **lon0** (*float*) – Longitude of position 0
- **lat0** (*float*) – Latitude of position 0
- **lon1** (*float*) – Longitude of position 1
- **lat1** (*float*) – Latitude of position 1

Returns

dist – Haversine distance between position 0 and position 1.

Return type

float

`GeoSpatialTools.distance_metrics.midpoint(lon0, lat0, lon1, lat1)`

Compute the midpoint of a great circle track

Parameters

- **lon0** (*float*) – Longitude of position 0
- **lat0** (*float*) – Latitude of position 0
- **lon1** (*float*) – Longitude of position 1
- **lat1** (*float*) – Latitude of position 1

Returns

Positions of midpoint between position 0 and position 1

Return type

lon, lat

9.3 Utils

Utility functions. Including Error classes and Warnings.

exception `GeoSpatialTools.utils.DateWarning`

Warning for Datetime Value

exception `GeoSpatialTools.utils.LatitudeError`

Error for invalid Latitude Value

PYTHON MODULE INDEX

g

`GeoSpatialTools.distance_metrics`, [27](#)
`GeoSpatialTools.great_circle`, [25](#)
`GeoSpatialTools.kdtree`, [22](#)
`GeoSpatialTools.neighbours`, [6](#)
`GeoSpatialTools.octtree`, [19](#)
`GeoSpatialTools.quadtree`, [14](#)
`GeoSpatialTools.record`, [7](#)
`GeoSpatialTools.shape`, [10](#)
`GeoSpatialTools.utils`, [28](#)

B

bearing() (in module *GeoSpatialTools.distance_metrics*), 27

C

cartesian_to_lonlat() (in module *GeoSpatialTools.great_circle*), 26

centre_datetime (*GeoSpatialTools.shape.SpaceTimeRectangle* property), 12

contains() (*GeoSpatialTools.shape.Ellipse* method), 10

contains() (*GeoSpatialTools.shape.Rectangle* method), 10

contains() (*GeoSpatialTools.shape.SpaceTimeEllipse* method), 11

contains() (*GeoSpatialTools.shape.SpaceTimeRectangle* method), 12

D

DateWarning, 28

delete() (*GeoSpatialTools.kdtree.KDTree* method), 22

destination() (in module *GeoSpatialTools.distance_metrics*), 27

dist_from_point() (*GeoSpatialTools.great_circle.GreatCircle* method), 25

distance() (*GeoSpatialTools.record.Record* method), 8

distance() (*GeoSpatialTools.record.SpaceTimeRecord* method), 8

divide() (*GeoSpatialTools.octtree.OctTree* method), 19

divide() (*GeoSpatialTools.quadtree.QuadTree* method), 15

E

edge_dist (*GeoSpatialTools.shape.Rectangle* property), 10

edge_dist (*GeoSpatialTools.shape.SpaceTimeRectangle* property), 12

Ellipse (class in *GeoSpatialTools.shape*), 10

F

find_nearest() (in module *GeoSpatialTools.neighbours*), 6

G

gcd_slc() (in module *GeoSpatialTools.distance_metrics*), 27

GeoSpatialTools.distance_metrics module, 27

GeoSpatialTools.great_circle module, 25

GeoSpatialTools.kdtree module, 22

GeoSpatialTools.neighbours module, 6

GeoSpatialTools.octtree module, 19

GeoSpatialTools.quadtree module, 14

GeoSpatialTools.record module, 7

GeoSpatialTools.shape module, 10

GeoSpatialTools.utils module, 28

GreatCircle (class in *GeoSpatialTools.great_circle*), 25

H

haversine() (in module *GeoSpatialTools.distance_metrics*), 28

I

insert() (*GeoSpatialTools.kdtree.KDTree* method), 22

insert() (*GeoSpatialTools.octtree.OctTree* method), 19

insert() (*GeoSpatialTools.quadtree.QuadTree* method), 15

intersection() (*GeoSpatialTools.great_circle.GreatCircle* method), 25

intersection_angle() (*GeoSpatialTools.great_circle.GreatCircle* method), 26

`intersects()` (*GeoSpatialTools.shape.Rectangle* method), 10
`intersects()` (*GeoSpatialTools.shape.SpaceTimeRectangle* method), 12

K

`KDTree` (class in *GeoSpatialTools.kdtree*), 22

L

`lat` (*GeoSpatialTools.shape.Rectangle* property), 10
`lat` (*GeoSpatialTools.shape.SpaceTimeRectangle* property), 12
`lat_range` (*GeoSpatialTools.shape.Rectangle* property), 11
`lat_range` (*GeoSpatialTools.shape.SpaceTimeRectangle* property), 12
`LatitudeError`, 28
`len()` (*GeoSpatialTools.octtree.OctTree* method), 19
`len()` (*GeoSpatialTools.quadtree.QuadTree* method), 15
`lon` (*GeoSpatialTools.shape.Rectangle* property), 11
`lon` (*GeoSpatialTools.shape.SpaceTimeRectangle* property), 12
`lon_range` (*GeoSpatialTools.shape.Rectangle* property), 11
`lon_range` (*GeoSpatialTools.shape.SpaceTimeRectangle* property), 12

M

`midpoint()` (in module *GeoSpatialTools.distance_metrics*), 28
module
 GeoSpatialTools.distance_metrics, 27
 GeoSpatialTools.great_circle, 25
 GeoSpatialTools.kdtree, 22
 GeoSpatialTools.neighbours, 6
 GeoSpatialTools.octtree, 19
 GeoSpatialTools.quadtree, 14
 GeoSpatialTools.record, 7
 GeoSpatialTools.shape, 10
 GeoSpatialTools.utils, 28

N

`nearby()` (*GeoSpatialTools.shape.Rectangle* method), 11
`nearby()` (*GeoSpatialTools.shape.SpaceTimeRectangle* method), 12
`nearby_points()` (*GeoSpatialTools.octtree.OctTree* method), 19
`nearby_points()` (*GeoSpatialTools.quadtree.QuadTree* method), 15

`nearby_rect()` (*GeoSpatialTools.shape.Ellipse* method), 10
`nearby_rect()` (*GeoSpatialTools.shape.SpaceTimeEllipse* method), 11

O

`OctTree` (class in *GeoSpatialTools.octtree*), 19

P

`polar_to_cartesian()` (in module *GeoSpatialTools.great_circle*), 26

Q

`QuadTree` (class in *GeoSpatialTools.quadtree*), 14
`query()` (*GeoSpatialTools.kdtree.KDTree* method), 23
`query()` (*GeoSpatialTools.octtree.OctTree* method), 20
`query()` (*GeoSpatialTools.quadtree.QuadTree* method), 15
`query_ellipse()` (*GeoSpatialTools.octtree.OctTree* method), 20
`query_ellipse()` (*GeoSpatialTools.quadtree.QuadTree* method), 15

R

`Record` (class in *GeoSpatialTools.record*), 7
`Rectangle` (class in *GeoSpatialTools.shape*), 10
`remove()` (*GeoSpatialTools.octtree.OctTree* method), 20
`remove()` (*GeoSpatialTools.quadtree.QuadTree* method), 16

S

`SortedError`, 6
`SortedWarning`, 6
`SpaceTimeEllipse` (class in *GeoSpatialTools.shape*), 11
`SpaceTimeRecord` (class in *GeoSpatialTools.record*), 8
`SpaceTimeRectangle` (class in *GeoSpatialTools.shape*), 11

T

`time_range` (*GeoSpatialTools.shape.SpaceTimeRectangle* property), 12